

```

hadoop@ritadmin:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [ritadmin]
Starting resourcemanager
Starting nodemanagers
hadoop@ritadmin:~$ jps
4626 SecondaryNameNode
4916 ResourceManager
5269 NodeManager
4376 DataNode
4090 NameNode
5550 Jps
hive> create database std;
OK
Time taken: 0.026 seconds
hive> use std
> ;
OK
Time taken: 0.011 seconds
hive> CREATE TABLE student (
>
>   id INT,
>
>   name STRING,
>
>   age INT,
>
>   department STRING,
>
>   grade STRING
> )
>
> ROW FORMAT DELIMITED
>
> FIELDS TERMINATED BY ','
>
> STORED AS TEXTFILE;
OK
Time taken: 0.026 seconds
hive> desc student;
OK
id            int
name          string
age           int
department    string
grade         string
Time taken: 0.019 seconds, Fetched: 5 row(s)

```

```
hive> LOAD DATA INPATH '/user/hive/warehouse/student_data/students.txt' INTO TABLE student;
```

Loading data to table std.student

OK

Time taken: 0.084 seconds

```
hive> select * from student;
```

OK

1	Don	10	CY	A
2	Alien	11	ISE	B
3	Bob	12	ECE	A
4	Rome	13	ME	B
5	Wick	15	CSE	A

Time taken: 0.051 seconds, Fetched: 5 row(s)

```
hive> create database emp;
```

OK

Time taken: 0.021 seconds

```
hive> use emp;
```

OK

Time taken: 0.012 seconds

```
hive> CREATE TABLE student (
```

```
> id INT,
```

```
>
```

```
> name STRING,
```

```
>
```

```
> age INT,
```

```
>
```

```
> department STRING,
```

```
>
```

```
> grade STRING
```

```
>
```

```
> )
```

```
>
```

```
> ROW FORMAT DELIMITED
```

```
>
```

```
> FIELDS TERMINATED BY '\t'
```

```
>
```

```
> STORED AS TEXTFILE;
```

OK

Time taken: 0.031 seconds

```
hive> LOAD DATA INPATH '/user/hive/warehouse/student_data/students.txt'
```

```
> INTO TABLE student;
```

Loading data to table emp.student

OK

Time taken: 0.076 seconds

```
hive> select * from student;
```

OK

1	Don	10	CY	A
2	Alien	11	ISE	B
3	Bob	12	ECE	A
4	Rome	13	ME	B
5	Wick	15	CSE	A

Time taken: 0.049 seconds, Fetched: 5 row(s)

hive>

```
hadoop@ritadmin: ~
OK
id                int
name              string
age               int
department         string
grade             string
Time taken: 0.019 seconds, Fetched: 5 row(s)
hive> LOAD DATA INPATH '/user/hive/warehouse/student_data/students.txt' INTO TABLE student;
FAILED: SemanticException line 1:17 Invalid path ''/user/hive/warehouse/student_data/students.txt'': No files matching path hdfs://localhost:9000/user/hive/warehouse/student_data/students.txt
hive> LOAD DATA INPATH '/user/hive/warehouse/student_data/students.txt' INTO TABLE student;
loading data to table std.student
OK
Time taken: 0.004 seconds
hive> select * from student;
OK
1      Don    10    CY    A
2      Allen  11    ISE  B
3      Bob    12    ECE  A
4      Rome   13    ME   B
5      Wick   15    CSE  A
Time taken: 0.049 seconds, Fetched: 5 row(s)
hive> select * from student;
OK
1      Don    10    CY    A
2      Allen  11    ISE  B
3      Bob    12    ECE  A
4      Rome   13    ME   B
5      Wick   15    CSE  A
Time taken: 0.051 seconds, Fetched: 5 row(s)
hive> create database emp;
OK
Time taken: 0.021 seconds
hive> use database emp;
NotableAltException(648[])
    at org.apache.hadoop.hive ql.parse.HiveParser_IdentifierParser_Identifier(HiveParser_IdentifierParser.java:11910)
    at org.apache.hadoop.hive ql.parse.HiveParser_Identifier(HiveParser.java:45972)
    at org.apache.hadoop.hive ql.parse.HiveParser_switchDatabaseStatement(HiveParser.java:6003)
    at org.apache.hadoop.hive ql.parse.HiveParser_ddlStatement(HiveParser.java:4207)
    at org.apache.hadoop.hive ql.parse.HiveParser_execStatement(HiveParser.java:2494)
    at org.apache.hadoop.hive ql.parse.HiveParser_statement(HiveParser.java:1420)
    at org.apache.hadoop.hive ql.parse.ParserDriver.parse(ParserDriver.java:220)
    at org.apache.hadoop.hive ql.parse.ParseUtils.parse(ParseUtils.java:74)
    at org.apache.hadoop.hive ql.parse.ParseUtils.parse(ParseUtils.java:67)
    at org.apache.hadoop.hive ql.Driver.compile(Driver.java:610)
    at org.apache.hadoop.hive ql.Driver.compileInternal(Driver.java:1826)
    at org.apache.hadoop.hive ql.Driver.compileAndRespond(Driver.java:1773)
    at org.apache.hadoop.hive ql.Driver.compileAndRespond(Driver.java:1768)
    at org.apache.hadoop.hive ql.rexec.ReExecDriver.compileAndRespond(ReExecDriver.java:120)
    at org.apache.hadoop.hive ql.rexec.ReExecDriver.run(ReExecDriver.java:214)
    at org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:239)
    at org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:188)
    at org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:402)
    at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:821)
    at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:759)
    at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:683)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:328)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:241)
FAILED: ParseException line 1:4 cannot recognize input near 'database' 'emp' '<EOF>' in switch database statement
hive> use emp;
OK
Time taken: 0.012 seconds
hive> CREATE TABLE student (
>   id INT,
>   name STRING,
>   age INT,
>   department STRING,
>   grade STRING
> )
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY '\t'
> STORED AS TEXTFILE;
OK
Time taken: 0.031 seconds
hive> LOAD DATA INPATH '/user/hive/warehouse/student_data/students.txt'
> INTO TABLE student;
FAILED: ParseException line 3:0 missing EOF at 'FIELDS' near 'student'
hive> LOAD DATA INPATH '/user/hive/warehouse/student_data/students.txt'
> INTO TABLE student;
FAILED: SemanticException line 1:17 Invalid path ''/user/hive/warehouse/student_data/students.txt'': No files matching path hdfs://localhost:9000/user/hive/warehouse/student_data/students.txt
hive> LOAD DATA INPATH '/user/hive/warehouse/student_data/students.txt'
> INTO TABLE student;
loading data to table emp.student
OK
Time taken: 0.070 seconds
hive> select * from student;
OK
1      Don    10    CY    A
2      Allen  11    ISE  B
3      Bob    12    ECE  A
4      Rome   13    ME   B
5      Wick   15    CSE  A
Time taken: 0.049 seconds, Fetched: 5 row(s)
hive>
```

```

hadoop@ritadmin: ~
Stage-Stage-3: Map: 1 Cumulative CPU: 1.31 sec HDFS Read: 9481 HDFS Write: 166 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 310 msec
OK
Alice Engineering
Bob HR
Charlie Marketing
Time taken: 14.214 seconds, Fetched: 3 row(s)
hive> select * from employees;
OK
1 Alice Engineering 70000.0
2 Bob HR 50000.0
3 Charlie Marketing 60000.0
Time taken: 0.053 seconds, Fetched: 3 row(s)
hive> select department, AVG(salary) as avg_salary
> from employees
> group by department;
Query ID = hadoop_20250407121049_2cb0a220-4e2b-4049-91f4-46939b401e38
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
set mapreduce.job.reduces=<number>
Starting Job = job_1744005325492_0006, Tracking URL = http://ritadmin:8088/proxy/applicatio
n_1744005325492_0006/
Kill Command = /home/hadoop/hadoop/bin/mapred job -kill job_1744005325492_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2025-04-07 12:10:55,099 Stage-1 map = 0%, reduce = 0%
2025-04-07 12:10:58,144 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.55 sec
2025-04-07 12:11:02,194 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.95 sec
MapReduce Total cumulative CPU time: 2 seconds 950 msec
Ended Job = job_1744005325492_0006
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.95 sec HDFS Read: 15645 HDFS Write:
172 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 950 msec
OK
Engineering 70000.0
HR 50000.0
Marketing 60000.0
Time taken: 14.002 seconds, Fetched: 3 row(s)
hive>

```

```

hadoop@ritadmin: ~
Ended Job = job_1744005325492_0004
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 1.31 sec HDFS Read: 9766 HDFS Write: 87 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 310 msec
OK
Time taken: 12.494 seconds
hive> select e.name, d.dept_name
> from employees e
> join departments d on (e.department=d.dept_name);
Query ID = hadoop_20250407120614_046ad41d-5d70-4934-a7ac-2ee81ed58ff8
Total jobs = 1
2025-04-07 12:06:17 Dump the side-table for tag: 1 with group count: 3 into file: file:
/tmp/hadoop/6f758721-de32-4b54-a969-8a6f3c1ec871/hive_2025-04-07_12-06-14_450_3961097237014
970895-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile11--.hashtable
2025-04-07 12:06:17 Uploaded 1 File to: file:/tmp/hadoop/6f758721-de32-4b54-a969-8a6f3c
1ec871/hive_2025-04-07_12-06-14_450_3961097237014970895-1/-local-10004/HashTable-Stage-3/Ma
pJoin-mapfile11--.hashtable (336 bytes)
2025-04-07 12:06:17 End of local task; Time Taken: 0.648 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1744005325492_0005, Tracking URL = http://ritadmin:8088/proxy/applicatio
n_1744005325492_0005/
Kill Command = /home/hadoop/hadoop/bin/mapred job -kill job_1744005325492_0005
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2025-04-07 12:06:23,564 Stage-3 map = 0%, reduce = 0%
2025-04-07 12:06:27,625 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.31 sec
MapReduce Total cumulative CPU time: 1 seconds 310 msec
Ended Job = job_1744005325492_0005
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 1.31 sec HDFS Read: 9481 HDFS Write: 166 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 310 msec
OK
Alice Engineering
Bob HR
Charlie Marketing
Time taken: 14.214 seconds, Fetched: 3 row(s)
hive> select * from employees;
OK
1 Alice Engineering 70000.0
2 Bob HR 50000.0
3 Charlie Marketing 60000.0
Time taken: 0.053 seconds, Fetched: 3 row(s)
hive>

```

```

hadoop@ritadmin: ~
n_1744005325492_0004/
Kill Command = /home/hadoop/hadoop/bin/mapred job -kill job_1744005325492_0004
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2025-04-07 12:05:38,069 Stage-3 map = 0%, reduce = 0%
2025-04-07 12:05:41,111 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.31 sec
MapReduce Total cumulative CPU time: 1 seconds 310 msec
Ended Job = job_1744005325492_0004
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 1.31 sec HDFS Read: 9766 HDFS Write: 87 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 310 msec
OK
Time taken: 12.494 seconds
hive> select e.name, d.dept_name
> from employees e
> join departments d on (e.department=d.dept_name);
Query ID = hadoop_20250407120614_046ad41d-5d70-4934-a7ac-2ee81ed58ff8
Total jobs = 1
2025-04-07 12:06:17 Dump the side-table for tag: 1 with group count: 3 into file: file:
/tmp/hadoop/6f758721-de32-4b54-a969-8a6f3c1ec871/hive_2025-04-07_12-06-14_450_3961097237014
970895-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile11--.hashtable
2025-04-07 12:06:17 Uploaded 1 File to: file:/tmp/hadoop/6f758721-de32-4b54-a969-8a6f3c
1ec871/hive_2025-04-07_12-06-14_450_3961097237014970895-1/-local-10004/HashTable-Stage-3/Ma
pJoin-mapfile11--.hashtable (336 bytes)
2025-04-07 12:06:17 End of local task; Time Taken: 0.648 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1744005325492_0005, Tracking URL = http://ritadmin:8088/proxy/applicatio
n_1744005325492_0005/
Kill Command = /home/hadoop/hadoop/bin/mapred job -kill job_1744005325492_0005
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2025-04-07 12:06:23,564 Stage-3 map = 0%, reduce = 0%
2025-04-07 12:06:27,625 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.31 sec
MapReduce Total cumulative CPU time: 1 seconds 310 msec
Ended Job = job_1744005325492_0005
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 1.31 sec HDFS Read: 9481 HDFS Write: 166 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 310 msec
OK
Alice Engineering
Bob HR
Charlie Marketing
Time taken: 14.214 seconds, Fetched: 3 row(s)
hive>

```

```

hadoop@ritadmin: ~
hadoop@ritadmin:~$ hive
Hive Session ID = 6f758721-de32-4b54-a969-8a6f3c1ec871

Logging initialized using configuration in jar:file:/home/hadoop/hive/lib/hive-common-3.1.2
.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consid
er using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Hive Session ID = 74817032-b578-4a35-94b9-fe0711956935
hive> use rit
> ;
OK
Time taken: 0.167 seconds
hive> SHOW TABLES;
OK
employees
Time taken: 0.078 seconds, Fetched: 1 row(s)
hive> SHOW TABLES LIKE 'department';
OK
Time taken: 0.032 seconds
hive> CREATE TABLE departments(
> id INT,
> dept_name STRING
> )
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> STORED AS TEXTFILE;
OK
Time taken: 0.21 seconds
hive> INSERT INTO TABLE departments VALUES
> (1, 'Engineering'),
> (2, 'HR'),
> (3, 'Marketing');
Query ID = hadoop_20250407115827_2abfb02d-1e92-4472-a628-1cc11bc4c064
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1744005325492_0003, Tracking URL = http://ritadmin:8088/proxy/applicatio
n_1744005325492_0003/
Kill Command = /home/hadoop/hadoop/bin/mapred job -kill job_1744005325492_0003
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 1
2025-04-07 12:06:33,069 Stage-3 map = 0%, reduce = 0%
2025-04-07 12:06:36,069 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.31 sec
MapReduce Total cumulative CPU time: 1 seconds 310 msec
Ended Job = job_1744005325492_0003
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 1.31 sec HDFS Read: 9481 HDFS Write: 166 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 310 msec
OK

```

```

hadoop@ritadmin: ~
2025-04-07 12:30:40,628 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.46 sec
2025-04-07 12:30:44,683 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.66 sec
MapReduce Total cumulative CPU time: 2 seconds 660 msec
Ended Job = job_1744005325492_0008
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.66 sec HDFS Read: 15087 HDFS Write:
153 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 660 msec
OK
1      70000.0
1      50000.0
1      60000.0
Time taken: 12.572 seconds, Fetched: 3 row(s)
hive> select count(name)as no_employees,department, MAX(salary) max_salary
> from employees
> group by department;
Query ID = hadoop_20250407123103_e07f2dee-9521-467c-bd3b-6a1abf554e86
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1744005325492_0009, Tracking URL = http://ritadmin:8088/proxy/applicatio
n_1744005325492_0009/
Kill Command = /home/hadoop/hadoop/bin/mapred job -kill job_1744005325492_0009
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2025-04-07 12:31:08,286 Stage-1 map = 0%, reduce = 0%
2025-04-07 12:31:12,337 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.57 sec
2025-04-07 12:31:16,384 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.7 sec
MapReduce Total cumulative CPU time: 2 seconds 700 msec
Ended Job = job_1744005325492_0009
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.7 sec HDFS Read: 15162 HDFS Write: 1
78 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 700 msec
OK
1      Engineering      70000.0
1      HR                50000.0
1      Marketing         60000.0
Time taken: 13.986 seconds, Fetched: 3 row(s)
hive>

```

```

hadoop@ritadmin:~$ hive
Hive Session ID = 6f4c730f-df35-486a-8a76-238854cb4172

Logging initialized using configuration in jar:file:/home/hadoop/hive/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Hive Session ID = d43c154b-98c1-4cbc-baea-f141b32df7d9

```

```

hive> select dept, count(*) from students group by dept;
Query ID = hadoop_20250409095631_3965b498-5b61-4936-9a18-81ad4aa56627
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1744170803527_0007, Tracking URL = http://ritadmin:8088/proxy/application_1744170803527_0007/
Kill Command = /home/hadoop/hadoop/bin/mapred job -kill job_1744170803527_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2025-04-09 09:56:36,935 Stage-1 map = 0%, reduce = 0%
2025-04-09 09:56:39,979 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.52 sec
2025-04-09 09:56:44,032 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.7 sec
MapReduce Total cumulative CPU time: 2 seconds 700 msec
Ended Job = job_1744170803527_0007
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.7 sec HDFS Read: 13808 HDFS Write: 218 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 700 msec
OK
Computer Science      2
Electrical Engineering 1
Electronics           1
Mechanical Engineering 1

```



```

hive> select dept, avg(age) from students group by dept;
Query ID = hadoop_20250409095708_6542f34e-a44b-46a2-8b68-b008d2c84817
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1744170803527_0008, Tracking URL = http://ritadmin:8088/proxy/application_1744170803527_0008/
Kill Command = /home/hadoop/hadoop/bin/mapred job -kill job_1744170803527_0008
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2025-04-09 09:57:13,102 Stage-1 map = 0%, reduce = 0%
2025-04-09 09:57:16,143 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.58 sec
2025-04-09 09:57:20,188 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.02 sec
MapReduce Total cumulative CPU time: 3 seconds 20 msec
Ended Job = job_1744170803527_0008
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.02 sec HDFS Read: 15900 HDFS Write: 230 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 20 msec
OK
Computer Science      21.0
Electrical Engineering 21.0
Electronics          20.0
Mechanical Engineering 19.0
Time taken: 12.77 seconds, Fetched: 4 row(s)

```

